

Introduction

The third assignment of the course focuses on the implementation and results of two supervised learning models, a linear regressor and an artificial neural network (ANN) with one hidden layer. Both models are trained to fit a one dimensional input-output dataset, and their performance is evaluated using the mean-squared error metric. Additionally, using hyperparameter adjustment, we observe the results of model complexity on learning ability.

Implementation Results

Part A

Q) Is it sufficient to use a linear regressor or is it necessary to use an ANN with a single hidden layer? If it is latter, what will be the minimum number of hidden units?

A: As will be seen in the later results in this report, it is necessary to use an artificial neural network with a single layer rather than a linear regressor, due to the fact that the distribution of the data is non-linear and linear regressor models cannot catch non-linear data. However, artificial neural networks can learn from data and fit accordingly.

Q) What is a good value for the learning rate?

A: Setting the learning rate as 0.001 provided sufficient results.

Q) How to initialize the weights?

A: Weights were initialized using small random values from a normal distribution scaled by 0.1.

Q) How many epochs should you use? How to decide when to stop?

A: Since the regression task is computationally easy, we can experiment with large values of epochs. There are two choices for stopping. Either the training phase ends with reach the upper bound of epoch, or we set a threshold for error, such that when the error gets desirably small, the training phase ends itself. However, this mechanism (**early stopping**) is not implemented in our model.

Q) Does normalization affect the learning process for this application?

A: Yes, normalization significantly improves the learning process. Without normalization, the model failed to converge due to large gradients and unstable updates.

Part B

The table given below displays the details of the artificial neural network model that we have implemented for the regression task:

ANN Used	64 Hidden Layers
Learning Rate	0.001
Range of initial weights	Normal distribution scaled by 0.1
Number of Epochs	1000
When to Stop	When epochs end (no early stopping)
Is Normalization Used	Yes
Training Loss (averaged)	46.6088
Test Loss (averaged)	78.8576

The figures given below display the results of the artificial neural network with parameters provided above, compared with the linear regressor model, and the loss curves per epoch:

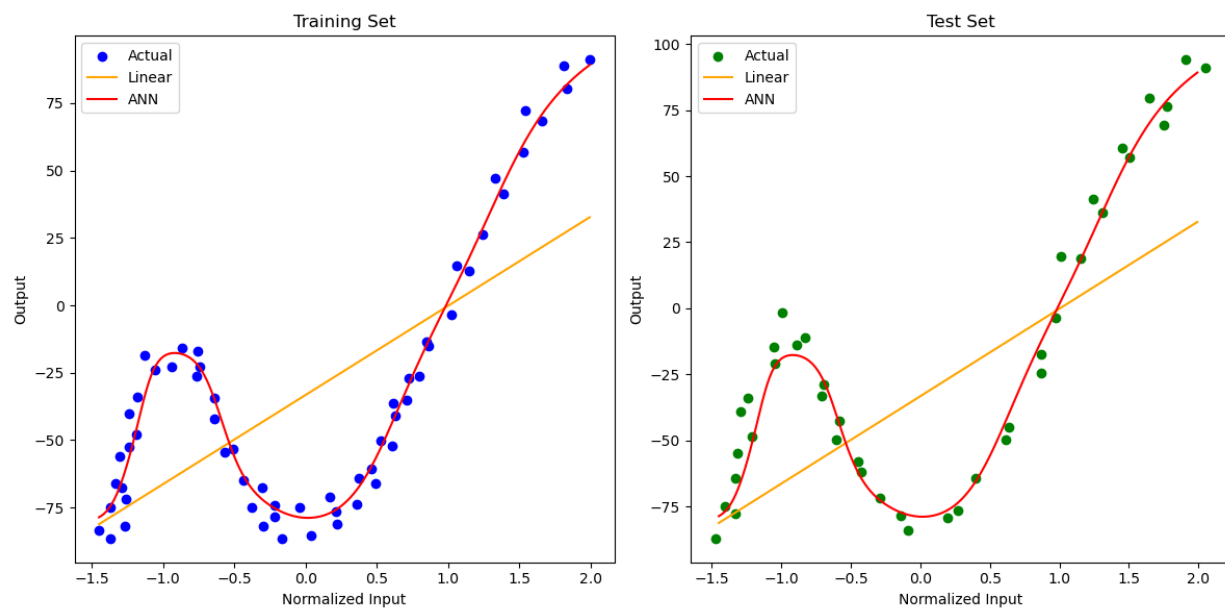


Figure 1: Results of the regression task on both training and test sets.




Figure 2: Training loss per epoch for both linear regressor and ANN regressor.

Part C

In this part, by using grid search, we have tested with multiple values of hyperparameters. The table given below displays the results of our tests with their respective hyperparameters:

Hidden Layers	Learning Rates	Epochs	Train MSE	Train STD	Test MSE	Test STD
2	0.01	1000	432.184	437.298	567.221	711.836
2	0.0005	1000	340.736	445.03	458.394	688.143
2	0.0001	1000	500.557	628.912	696.867	858.409
4	0.01	1000	350.489	479.92	497.099	744.636
4	0.0005	1000	332.782	452.1	448.083	698.826
4	0.0001	1000	445.564	532.514	615.564	767.223
8	0.01	1000	340.577	545.207	498.219	824.49
8	0.0005	1000	332.577	453.373	446.99	701.217
8	0.0001	1000	425.705	500.356	581.565	741.386
16	0.01	1000	355.668	607.332	530.888	902.679
16	0.0005	1000	333.684	454.212	447.215	702.938
16	0.0001	1000	417.586	481.648	560.524	727.018

32	0.01	1000	96.982	161.848	129.652	168.106
32	0.0005	1000	47.083	69.189	100.168	172.322
32	0.0001	1000	337.287	422.282	466.107	624.129

Inspecting the results of the parameter tuning yield that while increasing the number of neurons on the hidden layer increase the ability of capturing more complex features of the model, it may also lead to underfitting, hence it must be done with caution. On the other hand, learning rate adjustment is a crucial step in obtaining the optimal model. Decreasing the learning rate too much can also lead to underfitting, and increase computational costs, hence it must also be done with caution. The plots given below display results for various hidden layer numbers, and learning rates:




Figure 3: Results for 2 hidden layer neurons and a learning rate of 0.005.

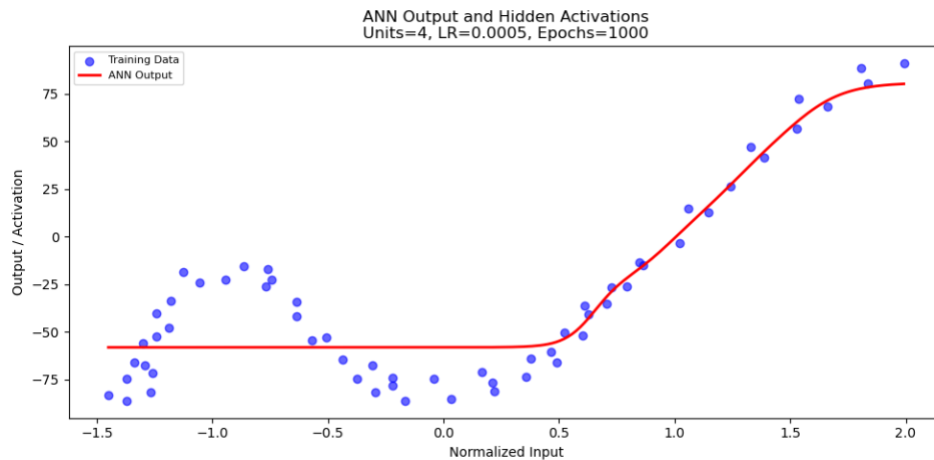


Figure 4: Results for 4 hidden layer neurons and a learning rate of 0.005.




Figure 5: Results for 8 hidden layer neurons and a learning rate of 0.005.




Figure 6: Results for 16 hidden layer neurons and a learning rate of 0.005.




Figure 7: Results for 32 hidden layer neurons and a learning rate of 0.005.